

Text, Screenshots, or Both?

Representation Value and Task Routing in Web Agents

Anonymous ACL submission

Abstract

A web agent can read a page as text, look at a screenshot, or do both, and deployments usually fix one choice for every task. We measure that choice across three sites of VisualWebArena and WebArena with four vision-language backbones, using text-only variants as controls and rerunning identical conditions to see which differences survive. Three findings separate questions that are usually merged. First, no observation choice is best everywhere, and the screenshot has a real, reproducible value that is predictable in advance on one site: a zero-cost rule over the task text picks out classifieds tasks on which adding the screenshot alone raises success by 25–30 points for the two strongest backbones (11 for a 4B model), and the gap reappears on a full rerun. Second, which representation suits which task is a genuine property of the task, beyond its difficulty, yet a single run measures it with reliability of only 0.31–0.50. Third, routers that choose a representation per task, with their operating point fixed on training data, do not convert this signal into a gain over the best fixed choices and their random mixtures. We describe the rerun-controlled protocol that keeps these three quantities apart.

1 Introduction

An agent that operates a website has to perceive it, and the two standard ways are complementary in principle: the page’s text (an accessibility tree) is precise and cheap, a screenshot shows what the text omits. A third option fuses them, overlaying numbered marks on the screenshot and listing the marked elements as text (Yang et al., 2023; Koh et al., 2024). Practitioners typically pick one of the three for a whole deployment. The question this paper asks is the one that choice raises: *is the choice worth making per task, and can it be made?*

Answering it requires keeping apart three quantities that a single leaderboard number merges:

1. **Representation value:** does the screenshot change success, holding everything else fixed?
2. **Preference reliability:** when one representation beats another on a task, is that a stable property of the task, or a property of one run?
3. **Routing gain:** can a policy that chooses per task, using only what is known before it runs, beat the fixed choices?

A large hindsight ceiling — pick, for each task, whichever representation happened to succeed — answers none of them, because it rewards run-to-run luck as much as structure. We therefore rerun identical conditions and read every comparison against what a rerun alone changes.

Findings. (i) No choice wins everywhere: the fused mode leads on classifieds, text leads on WebArena reddit, and the order reverses across backbones (§3). (ii) The screenshot has a specific value that can be isolated and predicted: with the text payload and prompt held fixed, adding the screenshot raises success on the classifieds tasks a zero-cost intent rule flags by 25–30 points for the two strongest backbones (11 for a 4B model), against 5–10 on the other tasks; the effect reproduces on a full rerun and is largest for the strongest backbone. On reddit the same rule finds nothing (§4). (iii) Per-task preferences between representations exist beyond task difficulty — a no-interaction null conditioned on difficulty is rejected in every fully rerun cell — but a single run measures them with reliability 0.31–0.50; two to three runs per task are needed to pass 0.5 (§5). (iv) Routers that fix their operating point on training data do not beat the fixed representations and their random mixtures; the room that exists above that frontier comes from task difficulty, not from representation fit, and is not captured without looking at test outcomes (§6).

Contributions. A rerun-controlled measurement of text, screenshot and fused observation across

three sites and four backbones, with text-only ablations that isolate the screenshot; a matched estimate of the screenshot’s value and an ex-ante rule that predicts it on one site; a test of whether per-task representation preferences are real, with a decision study for how many runs make them learnable; and a deployable-policy comparison against fixed choices *and their random mixtures* — the baseline of model-routing benchmarks, applied here to observation choice — which is the comparison a routing claim has to win (§7).

2 Setup

Sites and tasks. We use three sites of VisualWebArena (Koh et al., 2024) — classifieds (224 scored tasks), reddit (203) and shopping (432) — and the reddit site of WebArena (Zhou et al., 2024) (104 tasks, those all observation modes ran). VWA-reddit and WA-reddit are different task sets on the same application, so they are not two independent benchmarks. Scored sets are the collected tasks minus protocol exclusions written as rules over task configuration (Appendix A).

Backbones. B0 = Qwen3-VL-235B (API-served), B1 = Qwen3-VL-4B and B2 = Gemma-3-4B (both served locally), and B5 = GPT-5.6 (API, classifieds only). A *cell* is one site $\times$ backbone pair; the eight cells with every observation mode are VWA-classifieds and VWA-reddit $\times$ {B0, B1, B2} and WA-reddit $\times$ {B0, B1}. B5-classifieds and the two shopping cells extend them (§4, §6) and are never pooled with the eight.

One plain agent. One model call per step emits a thought, an action and its arguments; the prompt holds the instruction, the mode’s system prompt, the last eight steps and the current observation. There is no planner, memory, reflection or retry policy, so a difference between conditions is a difference of observation, not of scaffold. Episodes stop at 30 steps.

Observation modes. The deployment question has three answers: **DOM** (the page’s accessibility tree as text), **Vision** (the raw screenshot, with an empty text payload) and **SoM** (a set-of-marks screenshot plus the list of marked elements as text; Yang et al., 2023). Three further modes carry no screenshot and serve as controls; the one used throughout is **P-SoM**, which is SoM with the screenshot removed — same marks, same prompt — so SoM–P-SoM isolates the image (Appendix B




Figure 1: Success rate of the three deployment choices per cell. Shopping cells drop the tasks touched by a known harness defect in any mode (§2, Appendix D); B5’s Vision run is excluded for a coordinate-contract defect. Source: routing_three_arm, routing_extension_cells.

lists all six). “Text” modes withhold the *page* screenshot only: a reference image that is part of the task (“find this item”) is given in every mode.

Reruns. We reran identical conditions, 24 registered same-condition pairs in all; in three cells — VWA-classifieds B0, VWA-reddit B0 and WA-reddit B1 — every mode was run twice. A rerun uses the same tasks, code and site reset, so what it changes is what one run of the agent does not determine.

Scoring and cost. The evaluators are binary. Cost is the per-episode billed cost: an API invoice for B0 and B5, a token-priced electricity estimate for B1 and B2. The two kinds are ordered within a cell, never pooled across; §6 repeats the routing analysis under GPU-time and wall-clock cost.

3 No observation choice is best everywhere

Figure 1 shows the three deployment choices side by side. On classifieds the fused mode is the best of the three for B0, B1 and GPT-5.6 (B2 solves too few tasks to order the modes), and the margin over text is not small: SoM–DOM is +9.82 pp [+3.57, +16.07] for B0 and +8.04 pp [+3.57, +12.95] for B1 (paired task bootstrap). On WebArena reddit the order flips and text leads (SoM–DOM –4.81 pp for B0 and –2.88 pp for B1, both intervals crossing zero), and on VWA-reddit the weakest backbone does significantly worse with the fused mode (–2.96 pp [–5.91, –0.49]). On shopping the fused mode is best or tied with Vision (Figure 1).




Figure 2: Value of the screenshot alone (SoM–P-SoM: same marks, same prompt) on the tasks a zero-cost intent rule flags, against the other tasks. Classifieds is out-of-sample for the rule. “Rerun” is a full repeat of both modes. VWA-reddit B2 is omitted: no mode solves a flagged task. Source: visual_intent_routing.

Two cautions keep this descriptive. Cells on the same site share their tasks, so the eight are not eight independent replications. And the screenshot-free modes are four, not one: a maximum over four text variants against a single screenshot mode would be a selection effect, so the comparison here takes one representative per deployment class, as practitioners do.

What this section establishes is only that the choice matters and depends on the workload. Whether it can be made per task is the subject of the rest of the paper.

4 The screenshot has a value that can be isolated and predicted

A matched contrast. Comparing Vision with DOM changes two things at once — the image and the text the agent receives. SoM and P-SoM differ only in the annotated screenshot, so their difference is the value of the image alone. Chi et al. (2024) already compared screenshot-plus-text with text only on a subset of VisualWebArena tasks for one backbone (50.0% against 46.3%, without set-of-marks); we hold the set-of-marks text and prompt fixed, use four backbones and full task sets, rerun the contrast, and ask *where* the image pays off.

An ex-ante rule. A regular expression over the task intent flags tasks that ask about an image, a photo, a colour or a count of things shown, excluding tasks that already carry a reference image (Appendix C). It reads only the task configuration: no model call and no episode. It was written for

reddit, so its classifieds hits were never inspected and classifieds is out-of-sample.

Result. On the 71 classifieds tasks it flags, the screenshot raises success by +25.35 pp [+14.08, +36.62] for B0, against +5.23 pp on the other 153 tasks (Figure 2). The effect reproduces on a full rerun of both modes (+26.76 pp [+14.08, +39.44], from 23 against 4 successes), and it grows with capability: +2.82 pp for B2, +11.27 pp for B1, +25.35 pp for B0 and +29.58 pp [+16.90, +42.25] for GPT-5.6. On VWA-reddit the same rule flags 63 tasks and separates nothing: +3.17 pp on flagged tasks against +4.29 pp elsewhere for B0, +1.59 pp against −3.57 pp on its rerun. On shopping the flagged tasks carry at most five successes per mode, too few to read, and on WebArena the rule flags only 5 of 104 tasks.

Reading. Where the screenshot pays off on classifieds is predictable before anything runs, and it pays off more for stronger models; on reddit it is not. The screenshot’s value is therefore real and partly foreseeable, but site-specific: it is a statement about a kind of task, not a general law. It is also not a router. On these tasks the fused mode already contains the screenshot, so switching between DOM and Vision by this rule cannot beat simply using SoM; what remains open is whether *individual* tasks have stable preferences beyond such task types, which is the question of §5.

5 Per-task preferences are real, but one run cannot measure them

§4 found structure at the level of task *types*. Routing per task needs more: that an individual task prefers one representation over another, beyond being easy or hard, and that this preference is stable enough to learn from.

A null without task-by-mode structure. In the three cells where every mode was run twice, let $Y_{imr} \in \{0, 1\}$ be the outcome of task i under mode m in run r . Under

$$\text{logit Pr}(Y_{imr} = 1) = a_i + b_{mr},$$

every task has its own difficulty and every (mode, run) column its own easiness, but there is no persistent task-by-mode interaction. This is the Rasch structure that item-response routers use to model difficulty (Song et al., 2025); here it is the null, and the question is what lies beyond it. The row




Figure 3: Decision study for the three deployment modes: reliability of a task-by-mode label averaged over k runs, $\sigma_{\text{int}}^2 / (\sigma_{\text{int}}^2 + \sigma_{\text{err}}^2 / k)$. A single run gives 0.31–0.50. Source: routing_three_arm.

and column sums are sufficient for this model, so conditional on them every 0/1 matrix with those margins is equally likely; we sample that null exactly with curveball swaps, fitting nothing. The statistic is the reproducible part of the interaction, $\hat{\sigma}_{\text{int}}^2 = (\text{MS}_{\text{int}} - \text{MS}_{\text{err}}) / 2$ from the two-way task $\times$ mode ANOVA with two runs: the run-to-run covariance of each task’s mode profile.

Result. Over the three deployment modes the interaction exceeds the null in all three cells (classifieds B0 $\hat{\sigma}_{\text{int}}^2 = 0.033$ against a null 95th percentile of 0.015; reddit B0 0.018 against 0.009; WebArena-reddit B1 0.030 against 0.014; $p \leq 0.002$ each, Holm across cells). Over all six modes the same holds (Appendix E). Preferences between representations are therefore a property of tasks, not only of their difficulty.

But a single run measures them poorly. The same decomposition gives the reliability of a label built from k runs (Figure 3). From one run it is 0.34, 0.31 and 0.50 in the three cells; two to three runs per task are needed to pass 0.5. Over six modes it is lower still, 0.14–0.46, and up to six runs are needed. The preferences also persist where they can be observed directly: taking, for each task, the cheapest mode that succeeded on one run and scoring that choice on the rerun beats the fixed-mode frontier by 2.4–10.5 points on reddit B0 and WebArena-reddit B1. That choice uses the task’s own label, which no deployment has, so it measures persistence, not an attainable policy.

The routing target thus exists, and a training set built from single runs labels it with a reliability between a third and a half.

6 Deployable routers do not convert the signal

The comparison a router must win. A router that sends some tasks to an expensive mode and others to a cheap one trades success for cost. Beating every fixed mode on both axes is too strict, but beating the fixed modes one at a time is too lenient: sending a random share of tasks to each of two modes is also a deployable policy. We therefore compare against the *frontier of the fixed modes and all their random mixtures* — the upper concave envelope of the fixed points in the (cost, success) plane, the non-decreasing convex hull that model-routing benchmarks use as a baseline (Hu et al., 2024; Jitkrittum et al., 2025) — and report how far above it a router reaches.

Routers. Two constructions use 18 features available before a task runs: 14 intent-keyword indicators, the intent length, and three statistics of the first page. *Per-arm heads* fit one logistic success model per mode and send a task to the cheapest mode whose predicted success clears a threshold. *Triage* predicts whether any mode can solve the task, sends predicted-hopeless tasks to the cheapest mode and the rest to the best one. All predictions are out-of-fold (task-held-out five-fold).

Room above the frontier exists, and it is difficulty. Sweeping each router’s threshold traces a curve. Over the three deployment modes, the triage curve pooled over the eight cells rises 1.69 points above the frontier at low budget, which a label-shuffle null that refits the whole curve does not produce ($p = 0.001$; Table 1). The per-arm heads reach 0.86 points ($p = 0.072$). Triage does not choose a representation for a task; it decides which tasks are not worth the expensive mode. The room that exists is the difficulty lever, not representation fit.

A deployable router does not capture it. A curve’s best point is chosen after seeing test outcomes. When the operating point is instead chosen on training folds only, the pooled gain is +0.23 points for per-arm heads and +0.59 for triage, and neither interval excludes zero (Table 1). Over all six modes the deployable per-arm router sits significantly *below* the frontier (−0.38; 95% upper bound −0.14). The pattern holds for the strongest backbone: for GPT-5.6 on classifieds, where the screenshot’s value was largest (§4), the deployable

cells	policy	curve max (p)	deployable [5%, 95%]
3 arms, 8 cells pooled	per-arm heads	+0.86 (0.072)	+0.23 [−2.20, +0.22]
	triage	+1.69 (0.001)	+0.59 [−1.53, +1.23]
6 arms, 8 cells pooled	per-arm heads	+0.67 (0.252)	−0.38 [−3.20, −0.14]
	triage	+1.24 (0.043)	−0.39 [−2.42, +0.81]
GPT-5.6, classifieds (5 arms)	per-arm heads	+0.45 (0.768)	−2.45 [−8.61, +4.24]
	triage	+3.86 (0.127)	−1.63 [−6.70, +5.74]
Shopping B0 (3 arms)	per-arm heads	+1.14 (0.062)	−0.18 [−2.45, +2.02]
	triage	+0.20 (0.496)	−0.87 [−2.63, +1.37]
Shopping B1 (6 arms)	per-arm heads	+1.11 (0.143)	+0.56 [−4.69, +0.88]
	triage	+1.63 (0.158)	+1.19 [−4.10, +1.57]

Table 1: Routing gain over the fixed arms and their random mixtures, in SR points. *Curve max*: the best point of an out-of-fold curve, chosen after seeing test outcomes, with its label-shuffle p (pooled rows: max over a normalised budget). *Deployable*: the operating point chosen on training folds only, task bootstrap. Sources: routing_three_arm, representation_routing_frontier, routing_gain_upper_bounds, routing_extension_cells.

routers lose 1.6–2.5 points; on both shopping cells they stay within 1.2 points of the frontier.

Robustness. Holding tasks of the same template out together lowers the six-mode per-arm router’s mean deployable gain from −0.38 to −1.60 points, and fitting on one run and scoring on its rerun gives point estimates of both signs whose intervals all include zero (Appendix F). Repricing cost as GPU time or wall-clock time leaves the frontier result unchanged, but changes *which* single cell’s triage curve passes a per-cell Holm correction, so no single-cell result is quoted here. One of the three-mode intervals (per-arm heads, pooled) has an upper end below its point estimate, a sign that the bootstrap of a policy with a selected threshold is biased; we read these intervals as magnitudes, not as tests.

7 A protocol for multi-representation claims

The three quantities of §1 came apart because each was measured against its own control. We suggest the same four steps for any claim that choosing among several agent configurations per task would pay.

1. **Report reruns per arm and per cell**, not one flip rate or one resolution threshold for a whole study: in our data a rerun changes between 0 and 14% of outcomes depending on arm and cell (Appendix E).
2. **Isolate the component of interest with a matched ablation** (here SoM–P–SoM, which removes only the screenshot) before attributing a difference to it.

3. **Test for per-task structure against a null that conditions on difficulty**, and use the decision study to say how many runs per task a training label needs. A single-run oracle is not a ceiling a router can reach.

4. **Score a deployable policy, operating point chosen on training data, against the fixed arms and their random mixtures**, never only against the best single arm or a hindsight curve.

Applied here, the protocol says the screenshot is worth paying for on a predictable class of classifieds tasks, that per-task representation preferences are real but need two to three runs per task to be learnable, and that with single-run labels and pre-flight features a per-task representation router does not yet beat choosing one representation, or a fixed mixture, for the whole workload.

8 Related work

How web agents observe the page. Agents have read HTML or accessibility trees (Deng et al., 2023; Yang et al., 2025), screenshots grounded to elements (Zheng et al., 2024; He et al., 2024; Gou et al., 2025), or set-of-marks overlays that fuse the two (Yang et al., 2023; Koh et al., 2024). A parallel line shrinks or reshapes the text observation (Kerboua et al., 2025; Schiepanski and Piël, 2025; Enomoto et al., 2026), and some systems combine several grounding routes (Li et al., 2026a). Chi et al. (2024) compare screenshot-plus-text with text only on a VisualWebArena subset. Closer to adaptivity, Li and Shi (2026) gate visual observation components during execution of computer-use tasks, and Liu et al. (2026) route each action of a

computer-use agent among vision-language models; neither chooses between text and screenshot per task before it runs. These works compare observation designs at the level of a benchmark score or adapt them online; we ask whether the choice should be made per task in advance, and isolate the screenshot with a matched ablation rather than comparing whole pipelines.

Routing and cascades. Choosing among models per query to trade quality for cost is well studied for language models (Chen et al., 2024; Ong et al., 2025; Ding et al., 2024; Gupta et al., 2024; Peale et al., 2026); see Moslem and Kelleher (2026) for a survey. For web agents, Li et al. (2025) learn query-specific routing for cost-sensitive agents, and Hajimiri et al. (2026) judge agent components at a matched budget. Early exit (Lu et al., 2025) stops agents that are unlikely to succeed — the difficulty lever that, in our data, is the only part of routing that leaves room above the fixed frontier. Routing benchmarks compare routers against random mixtures of fixed models (Hu et al., 2024; Jitkrittum et al., 2025), and under unified conditions several learned routers fail to beat a simple baseline (Li et al., 2026b); Lee (2026) finds that assigning one model per task type in advance recovers most of the oracle routing gain on question answering, and that reruns rescore 5% of model–question pairs. Item-response routers model query difficulty and model ability (Song et al., 2025). We route over observation representations rather than models, take the mixture baseline from this literature, and add what it rarely has: reruns of identical agent conditions and a test of per-task structure beyond difficulty.

Reliability of agent evaluation. Repeated trials of the same agent disagree, which τ -bench measures with pass^k (Yao et al., 2025); inference itself is not deterministic even at fixed settings (He, 2025; Yuan et al., 2025), prompts change outcomes in ways unrelated to content (Sclar et al., 2024), and agent benchmarks carry evaluator and task defects (Hattami et al., 2025; Bhat et al., 2026). Our reruns quantify the first for web agents at the task level and show what it does to per-task routing labels; the harness defects we found are listed in Appendix D.

Limitations

Coverage of reruns. Only three cells have every mode run twice, two with the same backbone; the preference test and decision study (§5) rest on them.

Five of the eleven cells have no rerun of any mode.

Breadth. VWA-reddit and WA-reddit are task sets on the same application, so the study covers three applications, not four benchmarks. GPT-5.6 was run on classifieds only, and its Vision run is excluded for a coordinate-contract defect; a parser defect also let it take actions beyond the 30-step budget in every mode (Appendix D). The ex-ante rule of §4 works on classifieds only.

Shopping filtering. Dropping tasks touched by the search-box defect keeps the comparison between modes clean, but the dropped tasks are mostly search tasks, so the filtered cells read a different task mix; unfiltered results are reported alongside.

What the routing result does not show. The routers use 18 pre-flight features and logistic models; richer features, language-model routers or multiple runs per training task may do better. The result is that, with single-run labels and these features, the routers do not beat the fixed choices and their mixtures — not that per-task routing is impossible. Step-level switching is out of scope.

Exploratory analyses. The preference test, frontier comparisons and bounds were designed after the runs were collected. Each family states its multiplicity correction; the bootstrap intervals for policies with a selected threshold are read as magnitudes, not tests.

Cost. API invoices and electricity estimates are different quantities and are never pooled; local cost depends on the pricing basis, which is why §6 repeats the analysis under GPU time and wall-clock time.

References

- Vishvesh Bhat, Jay Vaghasiya, Muhammad Ahmed Mohsin, and Asad Aali. 2026. [Benchmarking the Benchmarks: A Validity Audit of Tool-Calling Evaluation](#). *Preprint*, arXiv:2607.02577.
- Lingjiao Chen, Matei Zaharia, and James Zou. 2024. [FrugalGPT: How to Use Large Language Models While Reducing Cost and Improving Performance](#). *Transactions on Machine Learning Research*.
- Wayne Chi, Ameet Talwalkar, and Chris Donahue. 2024. [The impact of element ordering on LM agent performance](#). *Preprint*, arXiv:2409.12089.

486	Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, Sam Stevens, Boshi Wang, Huan Sun, and Yu Su. 2023. Mind2Web: Towards a Generalist Agent for the Web . In <i>Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)</i> .	542
487		543
488		544
489		545
490		546
491	Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Rühle, Laks V. S. Lakshmanan, and Ahmed Hassan Awadallah. 2024. Hybrid LLM: Cost-Efficient and Quality-Aware Query Routing . In <i>International Conference on Learning Representations</i> .	547
492		548
493		549
494		550
495		551
496		552
497	Masafumi Enomoto, Ryoma Obara, Haochen Zhang, and Masafumi Oyamada. 2026. Revisiting Observation Reduction for Web Agents: Comprehensive Evaluation with a Lightweight Framework . <i>Preprint</i> , arXiv:2605.29397.	553
498		554
499		555
500		556
501		557
502	Boyu Gou, Demi Ruohan Wang, Boyuan Zheng, Yanan Xie, Cheng Chang, Yiheng Shu, Huan Sun, and Yu Su. 2025. Navigating the Digital World as Humans Do: Universal Visual Grounding for GUI Agents . In <i>International Conference on Learning Representations</i> .	558
503		559
504		
505		560
506		561
507		
508	Neha Gupta, Harikrishna Narasimhan, Wittawat Jitkritum, Ankit Singh Rawat, Aditya Krishna Menon, and Sanjiv Kumar. 2024. Language Model Cascades: Token-Level Uncertainty and Beyond . In <i>International Conference on Learning Representations</i> .	562
509		563
510		564
511		565
512		
513	Sina Hajimiri, Masih Aminbeidokhti, Jose Dolz, Ismail Ben Ayed, Issam H. Laradji, Spandana Gella, and Nicolas Gontier. 2026. Are Online Skill and Memory Modules Always Worth Their Tokens? A Budget-Constrained Study of Web Agents . <i>Preprint</i> , arXiv:2606.15017. Accepted to EMNLP 2026 according to author lab; proceedings not verified.	566
514		567
515		568
516		
517		569
518		570
519		571
520	Amine El Hattami, Megh Thakkar, Nicolas Chapados, and Christopher Pal. 2025. WebArena Verified: Reliable Evaluation for Web Agents . In <i>Workshop on Scaling Environments for Agents</i> . NeurIPS 2025 workshop; also see https://github.com/ServiceNow/webarena-verified .	572
521		573
522		574
523		575
524		
525		576
526	Hongliang He, Wenlin Yao, Kaixin Ma, Wenhao Yu, Yong Dai, Hongming Zhang, Zhenzhong Lan, and Dong Yu. 2024. WebVoyager: Building an End-to-End Web Agent with Large Multimodal Models . In <i>Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics</i> .	577
527		578
528		579
529		
530		580
531		581
532	Horace He. 2025. Defeating Nondeterminism in LLM Inference . Thinking Machines Lab technical blog; not peer-reviewed.	582
533		583
534		
535	Qitian Hu, Jacob Bieker, Xiuyu Li, Nan Jiang, Benjamin Keigwin, Gaurav Ranganath, Kurt Keutzer, and Shriyash Upadhyay. 2024. RouterBench: A benchmark for multi-LLM routing system . In <i>Agentic Markets Workshop at ICML 2024</i> .	584
536		585
537		586
538		587
539		588
540	Wittawat Jitkritum, Harikrishna Narasimhan, Ankit Singh Rawat, Jeevesh Juneja, Congchao	589
541		590
		591
		592
		593
		594
		595
	Wang, Zifeng Wang, Alec Go, Chen-Yu Lee, Pradeep Shenoy, Rina Panigrahy, Aditya Krishna Menon, and Sanjiv Kumar. 2025. Universal model routing for efficient LLM inference . <i>Preprint</i> , arXiv:2502.08773.	
	Imene Kerboua, Sahar Omid Shayegan, Xing Han Lù, Léo Boisvert, Megh Thakkar, Massimo Caccia, Jérémy Espinas, Alexandre Aussem, Véronique Eglin, and Alexandre Lacoste. 2025. FocusAgent: Simple yet effective ways of trimming the large context of web agents . <i>Preprint</i> , arXiv:2510.03204.	
	Jing Yu Koh, Robert Lo, Lawrence Jang, Vikram Duvvur, Ming Lim, Po-Yu Huang, Graham Neubig, Shuyan Zhou, Russ Salakhutdinov, and Daniel Fried. 2024. VisualWebArena: Evaluating Multimodal Agents on Realistic Visual Web Tasks . In <i>Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics</i> .	
	Janghoon Lee. 2026. Most of the LLM routing gap is task type . <i>Preprint</i> , arXiv:2608.23023.	
	Aiden Yiliu Li, Xinyue Hao, Shilong Liu, and Mengdi Wang. 2026a. Avenir-Web: Human-Experience-Imitating Multimodal Web Agents with Mixture of Grounding Experts . <i>Preprint</i> , arXiv:2602.02468.	
	Bojie Li and Noah Shi. 2026. Agent-computer observation interfaces enable dynamic computer use . <i>Preprint</i> , arXiv:2606.29472.	
	Hao Li, Yiqun Zhang, Zhaoyan Guo, Chenxu Wang, Shengji Tang, Qiaosheng Zhang, Yang Chen, Biqing Qi, Peng Ye, Lei Bai, Zhen Wang, and Shuyue Hu. 2026b. LLMRouterBench: A massive benchmark and unified framework for LLM routing . In <i>Findings of the Association for Computational Linguistics: ACL 2026</i> .	
	Tao Li, Jinlong Hu, Yang Wang, Junfeng Liu, and Xuejun Liu. 2025. WebRouter: Query-specific Router via Variational Information Bottleneck for Cost-sensitive Web Agent . <i>Preprint</i> , arXiv:2510.11221.	
	Xunzhuo Liu, Bowei He, Xue Liu, Andy Luo, Haichen Zhang, and Huamin Chen. 2026. Adaptive vision-language model routing for computer use agents . <i>Preprint</i> , arXiv:2603.12823.	
	Qingyu Lu, Liang Ding, Siyi Cao, Xuebo Liu, Kanjian Zhang, Jinxia Zhang, and Dacheng Tao. 2025. Run-away is Ashamed, But Helpful: On the Early-Exit Behavior of Large Language Model-based Agents in Embodied Environments . In <i>Findings of the Association for Computational Linguistics: EMNLP 2025</i> .	
	Yasmin Moslem and John D. Kelleher. 2026. Dynamic Model Routing and Cascading for Efficient LLM Inference: A Survey . <i>Transactions on Machine Learning Research</i> . Survey Certification; confirmed by author repository.	

Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez, M. Waleed Kadous, and Ion Stoica. 2025. [RouteLLM: Learning to Route LLMs from Preference Data](#). In *International Conference on Learning Representations*.

Charlotte Peale, Siddhartha Devic, Parikshit Gopalan, Udi Wieder, and Aravind Gollakota. 2026. [Flexible Routing via Uncertainty Decomposition](#). *Preprint*, arXiv:2605.07805. ICML 2026 Workshop on Statistical Frameworks for Uncertainty in Agentic Systems.

Thassilo M. Schiepaniski and Nicholas Piël. 2025. [Beyond Pixels: Exploring DOM Downsampling for LLM-Based Web Agents](#). *Preprint*, arXiv:2508.04412.

Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. 2024. [Quantifying Language Models’ Sensitivity to Spurious Features in Prompt Design or: How I learned to start worrying about prompt formatting](#). In *International Conference on Learning Representations*.

Wei Song, Zhenya Huang, Cheng Cheng, Weibo Gao, Bihan Xu, GuanHao Zhao, Fei Wang, and Runze Wu. 2025. [IRT-Router: Effective and interpretable multi-LLM routing via item response theory](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15629–15644.

Jianwei Yang, Hao Zhang, Feng Li, Xueyan Zou, Chunyuan Li, and Jianfeng Gao. 2023. [Set-of-Mark Prompting Unleashes Extraordinary Visual Grounding in GPT-4V](#). *Preprint*, arXiv:2310.11441.

Ke Yang, Yao Liu, Sapana Chaudhary, Rasool Fakoor, Pratik A. Chaudhari, George Karypis, and Huzefa Rangwala. 2025. [AgentOccam: A Simple Yet Strong Baseline for LLM-Based Web Agents](#). In *International Conference on Learning Representations*.

Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. 2025. [\$\tau\$ -bench: A benchmark for tool-agent-user interaction in real-world domains](#). In *International Conference on Learning Representations*.

Jiayi Yuan, Hao Li, Xinheng Ding, Wenya Xie, Yujie Li, Wentian Zhao, Kun Wan, Jing Shi, Xia Hu, and Zirui Liu. 2025. [Understanding and Mitigating Numerical Sources of Nondeterminism in LLM Inference](#). *Preprint*, arXiv:2506.09501.

Boyuan Zheng, Boyu Gou, Jihyung Kil, Huan Sun, and Yu Su. 2024. [GPT-4V\(ision\) is a Generalist Web Agent, if Grounded](#). In *Proceedings of the 41st International Conference on Machine Learning*.

Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. 2024. [WebArena: A Realistic Web Environment for Building Autonomous Agents](#). In *International Conference on Learning Representations*.

A Task universes and exclusions

Scored sets are the collected tasks minus protocol exclusions written as rules over task configuration: classifieds 224 of 224; VWA-reddit 203 of 205 (one task whose evaluator never checks the requested action, and one cross-site task answerable from memory); shopping 432 of 435; WebArena-reddit the 104 tasks every mode ran. Every analysis records the task set it was built on and fails if denominators differ.

B The six observation modes

Four modes carry no page screenshot and cross two factors: the text payload (the accessibility tree, or the flattened list of set-of-marks elements) and the prompt family (DOM or SoM). DOM is tree + DOM prompt; P-prompt is tree + SoM prompt; P-text is marks + DOM prompt; P-SoM is marks + SoM prompt. SoM adds the annotated screenshot to P-SoM; Vision is the raw screenshot with an empty text payload. A task’s own reference image, when it has one, is given in every mode.

C The intent rule

A task is flagged if its intent matches
`\b(image|picture|photo|screenshot)\b|`
`\bcolour?r of\b|`
`\bhow many\b[^\.]{0,40}\bin(?:the|this)\b`

(case-insensitive) and the task carries no reference image. It flags 71 of 224 classifieds tasks, 63 of 203 VWA-reddit tasks, 56 of 432 shopping tasks and 5 of 104 WebArena-reddit tasks. It was written while auditing reddit failures; classifieds is out-of-sample.

D Harness defects and how cells were filtered

- **GPT-5.6 Vision:** the run used a broken coordinate contract and is excluded.
- **GPT-5.6 step budget:** when the model emitted several valid actions, the parser executed one but booked the step as a free wait, so each GPT-5.6 condition took 50–90 actions beyond the 30-step budget (2–24 for B2, none for B0 and B1).
- **Shopping search box:** typing into a search box that already holds a query submitted old + new text (“chair” + “chair with wheels” → “chairchair with wheels”) in 43% of such submissions, far more often in text modes than

699 in Vision. A detector reproducing the defect's
700 per-mode counts marks affected episodes; the
701 *clean* shopping cells drop every task affected
702 in any mode (B0 keeps 368 of 432 tasks, B1
703 192).

- 704 • **Shopping state:** the category grid order dif-
705 fers between runs (42 order-sensitive tasks
706 dropped), and the wishlist is not reset between
707 tasks (4 tasks dropped).

708 **E Six-mode results**

709 Over all six modes, the no-interaction null of §5 is
710 rejected in all three fully rerun cells ($p = 0.0005$,
711 0.017, 0.0005; Holm), with single-run reliability
712 0.29, 0.14 and 0.46 and three, six and two runs per
713 task needed to reach 0.5. Across the 24 registered
714 rerun pairs, a rerun changes between 0 and 14% of
715 task outcomes, depending on arm and cell.

716 **F Template and rerun robustness**

717 Averaged over cells, the deployable excess over
718 the frontier with folds grouped by task template is
719 -1.60 (per-arm heads) and -0.84 (triage) points
720 over six modes, and -0.64 and $+0.29$ over three.
721 Fitting on one run and scoring on its rerun, aver-
722 aged over the three fully rerun cells, gives -0.90
723 and -1.60 over six modes and -0.60 and -0.56
724 over three. Over six modes, every per-cell template-
725 bootstrap interval (90%) includes zero; the three-
726 mode readings are point estimates.