



Testing when richer web-agent representations help — and whether their value can be predicted cheaply

WHAT AN AGENT SEES, AND WHO CHOOSES



THE SAME TASK, SEEN TWO WAYS


“Go to my listing of the blue bike and change the price to \$85.50 — and say so in the description.”



1 · SIX VIEWS, EIGHT SETTINGS


	DOM	P-prompt	P-text	P-SoM	SoM	Vision
no image	17.4	14.3	6.2	5.9	1.3	3.9
text + image	19.6	12.3	6.7	5.4	1.8	0.0
image only	15.6	13.3	7.6	5.9	0.5	2.0
no image	15.6	10.8	6.7	5.9	0.9	0.5
text + image	27.2	14.8	14.3	7.4	2.2	1.0
image only	25.0	7.4	12.5	2.5	2.2	2.0

2 · THEY SOLVE DIFFERENT TASKS




The sets overlap — but do not coincide.

3 · THEY BEHAVE DIFFERENTLY




Vision scrolls ~4x more.

4 · AND THEY FAIL DIFFERENTLY




One side dies of something you can name; the other never arrives.

5 · SO CHOOSE PER TASK? NOT SO FAST.



A win lands in the shaded region. None does.

6 · AND THIS IS WHY



More routing upside, less usable training signal.

