


Testing when richer web-agent representations help — and whether their value can be predicted cheaply

WHAT AN AGENT SEES, AND WHO CHOOSES



THE SAME TASK, SEEN TWO WAYS

“Go to my listing of the blue bike and change the price to \$85.50 — and say so in the description.”

READ page text only

12 steps
8 pages seen
\$0.09 solved

LOOK screenshot only


26 steps
9 pages seen
\$0.10 gave up

Both start on the same form. Then READ finishes it, while LOOK ends up back on the page it started from. One recorded run per view; both outcomes repeat on an independent rerun.

1 · SIX VIEWS, EIGHT SETTINGS


	DOM	P-prompt	P-text	P-SoM	SoM	Vision
cls-80 n=224	17.4	14.3	6.2	5.9	1.3	3.9
cls-81 n=203	19.4	12.3	6.7	5.4	1.8	0.0
cls-82 n=224	15.6	13.3	7.6	5.9	0.5	2.0
cls-83 n=203	15.6	10.8	6.7	5.9	0.9	0.5
cls-84 n=224	27.2	14.8	14.3	7.4	2.2	1.0
cls-85 n=203	25.0	7.4	12.5	2.5	2.2	2.0
cls-86 n=104						
cls-87 n=104						
cls-88 n=104						
cls-89 n=104						
cls-90 n=104						
cls-91 n=104						
cls-92 n=104						
cls-93 n=104						
cls-94 n=104						
cls-95 n=104						
cls-96 n=104						
cls-97 n=104						
cls-98 n=104						
cls-99 n=104						
cls-100 n=104						

2 · THEY SOLVE DIFFERENT TASKS




The sets overlap — but do not coincide.

3 · THEY BEHAVE DIFFERENTLY




Vision scrolls ~4x more.

4 · AND THEY FAIL DIFFERENTLY




Text-only skews toward early give-up; image-only toward stalled progress.

5 · SO CHOOSE PER TASK? NOT SO FAST.



No learned policy enters the win region; the lone square is hindsight.

6 · AND THIS IS WHY



More routing upside, less usable training signal.

