



Testing when richer web-agent representations help — and whether their value can be predicted cheaply

WHAT AN AGENT SEES, AND WHO CHOOSES




THE SAME TASK, SEEN TWO WAYS


“Go to my listing of the blue bike and change the price to \$85.50 — and say so in the description.”




1 · SIX VIEWS, EIGHT SETTINGS




2 · THEY SOLVE DIFFERENT TASKS



3 · THEY BEHAVE DIFFERENTLY




4 · AND THEY FAIL DIFFERENTLY




One side dies of something you can name; the other never arrives.

5 · SO CHOOSE PER TASK? NOT SO FAST.



A win lands in the shaded region. None does.

6 · AND THIS IS WHY



More routing upside, less usable training signal.

