



## Testing when richer web-agent representations help — and whether their value can be predicted cheaply

### WHAT AN AGENT SEES, AND WHO CHOOSES




### THE SAME TASK, SEEN TWO WAYS


"Go to my listing of the blue bike and change the price to \$85.50 — and say so in the description."



### 1 · SIX VIEWS, EIGHT SETTINGS




### 2 · THEY SOLVE DIFFERENT TASKS




The sets overlap — but do not coincide.

### 3 · THEY BEHAVE DIFFERENTLY




Vision scrolls ~4x more.

### 4 · AND THEY FAIL DIFFERENTLY




One side dies of something you can name; the other never arrives.

### 5 · SO CHOOSE PER TASK? NOT SO FAST.



A win lands in the shaded region. None does.

### 6 · AND THIS IS WHY



More routing upside, less usable training signal.

