

When is expensive perception worth paying for?

Jiaming Wei

Supervisors: Prof. María Pérez-Ortiz · Zekun Wu

MSc thesis · REALM workshop, EMNLP 2026

■ A web agent works the page for you — and can see it three ways.

**Codex**
computer use

Create account

Username
Avoid using your real name. ⓘ

Password

Confirm password

Email address (optional)

This site is protected by hCaptcha and its [Privacy Policy](#) and [Terms of Service](#) apply.

LOOK · screenshot

**Claude Code**
Playwright MCP

playwright · browser_snapshot — real reply 161 lines

```
- generic [active] [ref=e1]:  
: 26 lines  
- heading "Create account" [level=1] [ref=e35]  
: 55 lines  
- generic "Create account" [ref=e53]:  
: 6 lines  
  - button "username policy" [ref=e62]  
  - textbox "Username" [ref=e66]:  
    - /placeholder: Enter your username  
: 3 lines  
  - textbox "Password" [ref=e72]:  
    - /placeholder: Enter a password  
: 4 lines  
  - textbox "Confirm password" [ref=e79]:  
    - /placeholder: Enter password again  
: 3 lines  
  - textbox "Email address (optional)" [ref=e86]:  
    - /placeholder: Enter your email address  
  - button "Create your account" [ref=e89]  
: 51 more lines
```

READ · page as text

**browser-use**
highlighted elements

Create account

Username
Avoid using your real name. ⓘ

Password

Confirm password

Email address (optional)

This site is protected by hCaptcha and its [Privacy Policy](#) and [Terms of Service](#) apply.

BOTH · screenshot + marks

One real page, 15 Sep 2026 · cropped to the form · logos name the tools, not their limits

ONE RECORDED RUN PER VIEW · NOTHING ON THIS SCREEN IS LIVE

■ Same task, three ways of seeing it, three different bills.

■ Choose the right view for every task: 27 → 43 of 100 solved.



Upper bound · the demo's site, 224 tasks · other two sites +11 and +16 per 100, -14 and -27% · not Claude or GPT results

IF YOU BUILD A WEB AGENT

Should it look, read, or both — and can it learn to choose?

1

How do the views behave?


2

How do they fail?

3


Can it be learned?

■ With only the screenshot, the agent scrolls far more and types far less.



Medians over 8 settings: LOOK scrolls in **about 30%** of its steps, every other view **about 6%**; it types about half as often


- Text-only views fail in ways you can name; views with the screenshot just never get there.



Each side vs its own failures elsewhere · text-only: several over **2x** · LOOK and BOTH: none stands out · six VisualWebArena settings

③ CAN A ROUTER LEARN WHICH VIEW TO USE?

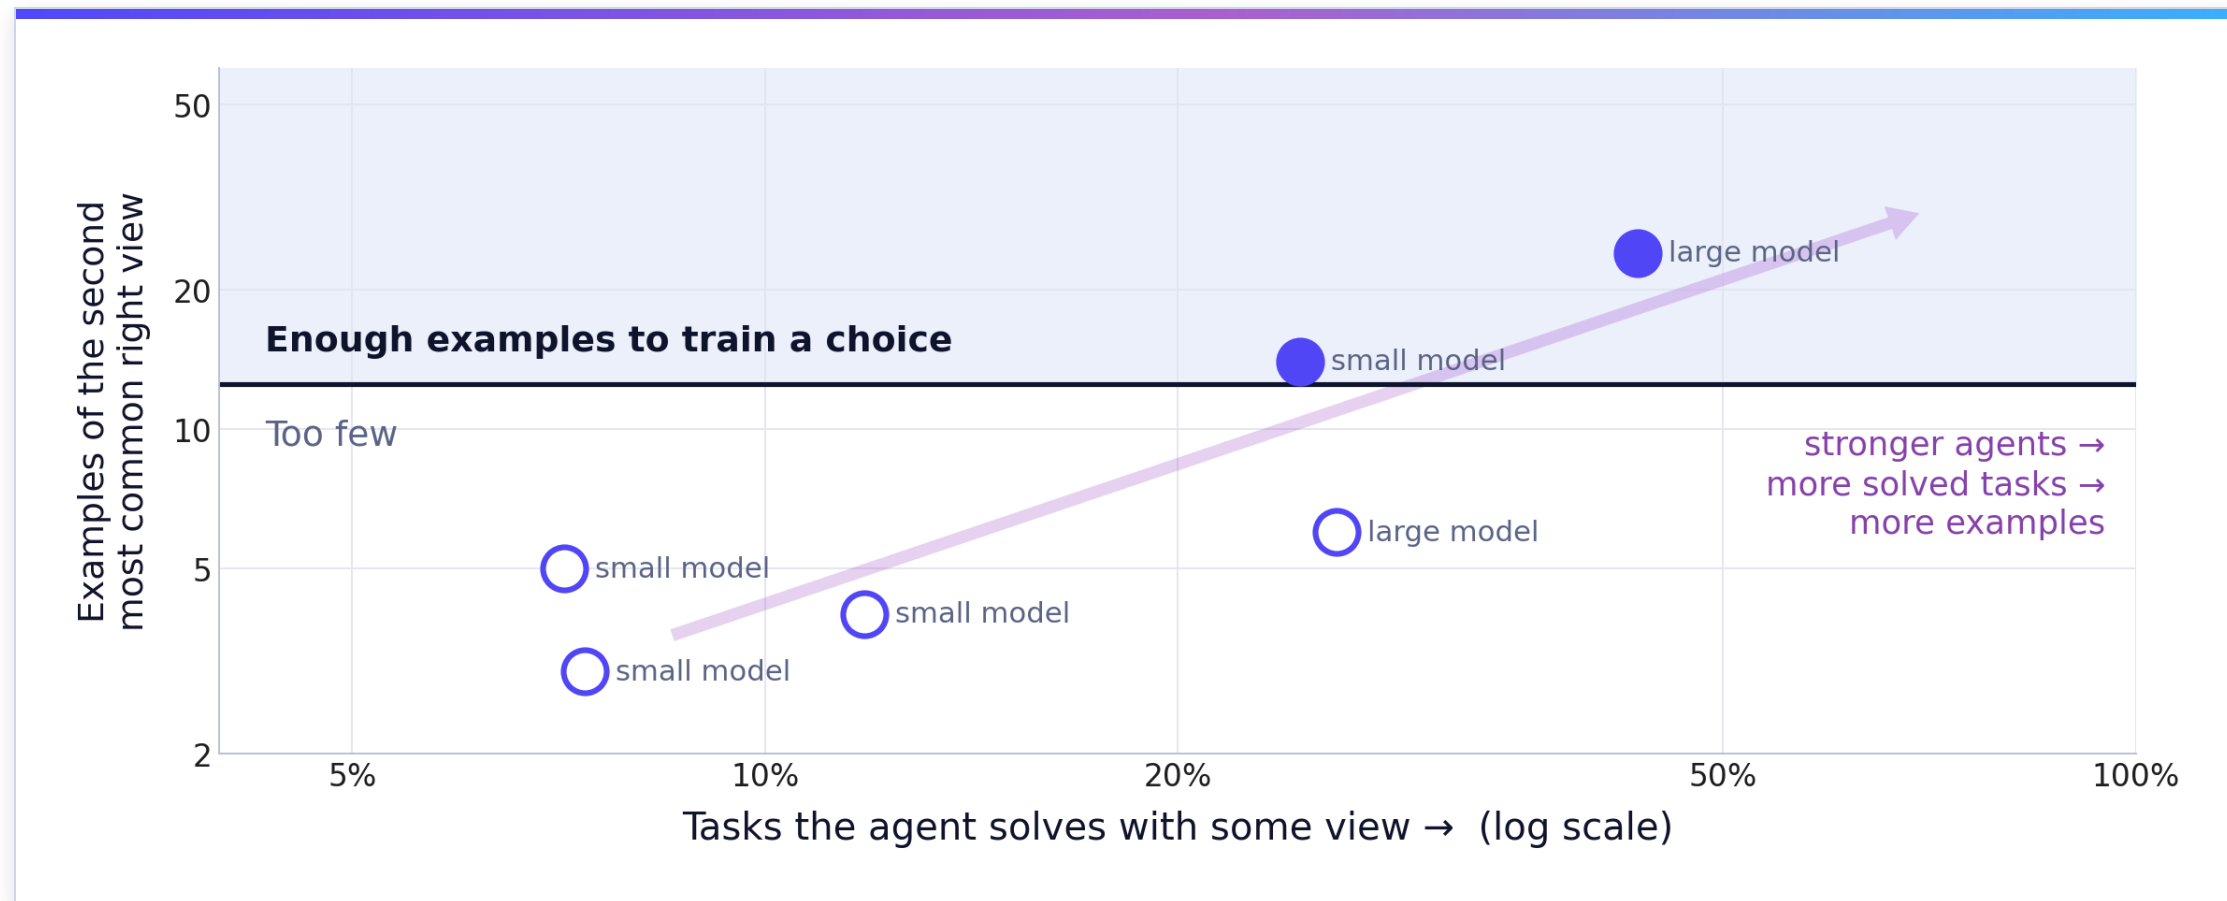
■ No learned router reaches the win corner — even a perfect one rarely does.



Learned routers: **0 of 8** settings

Perfect router: **1 of 8**

■ Routing's scaling law: no wins, no examples.



Six VisualWebArena settings · the arrow is a direction, not a fit · below the line, today's agents need at least **2–4×** more tasks

NOT YET

■ Improve the agent first, collect reliable examples, then learn when to look.

WHAT IT TAKES

- a stronger agent
- at least **2–4×** more tasks
- reruns that agree

THREE ASKS

- ✓ type a task at the board
- ✓ which view do you use?
- ✓ lend us a stronger agent

Look, read, or both?
It depends on the task.
Nothing we trained knows it yet.



Come to the board — try your own task, live in all three views.

Jiaming Wei · jiaming.wei.25@ucl.ac.uk

REFERENCE – TASK 130, RECORDED STEPPING, 39 S