

When is expensive perception worth paying for?

Jiaming Wei

Supervisors: Prof. María Pérez-Ortiz · Zekun Wu

MSc thesis · REALM workshop, EMNLP 2026

ONE RECORDED RUN PER VIEW · NOTHING ON THIS SCREEN IS LIVE

■ Same task, three ways of seeing it, three different bills.

IF YOU BUILD A WEB AGENT

Should it look, read, or both — and can it learn to choose?

1

How do the views
behave?

2

How do they fail?


3

What does picking
right buy?

4

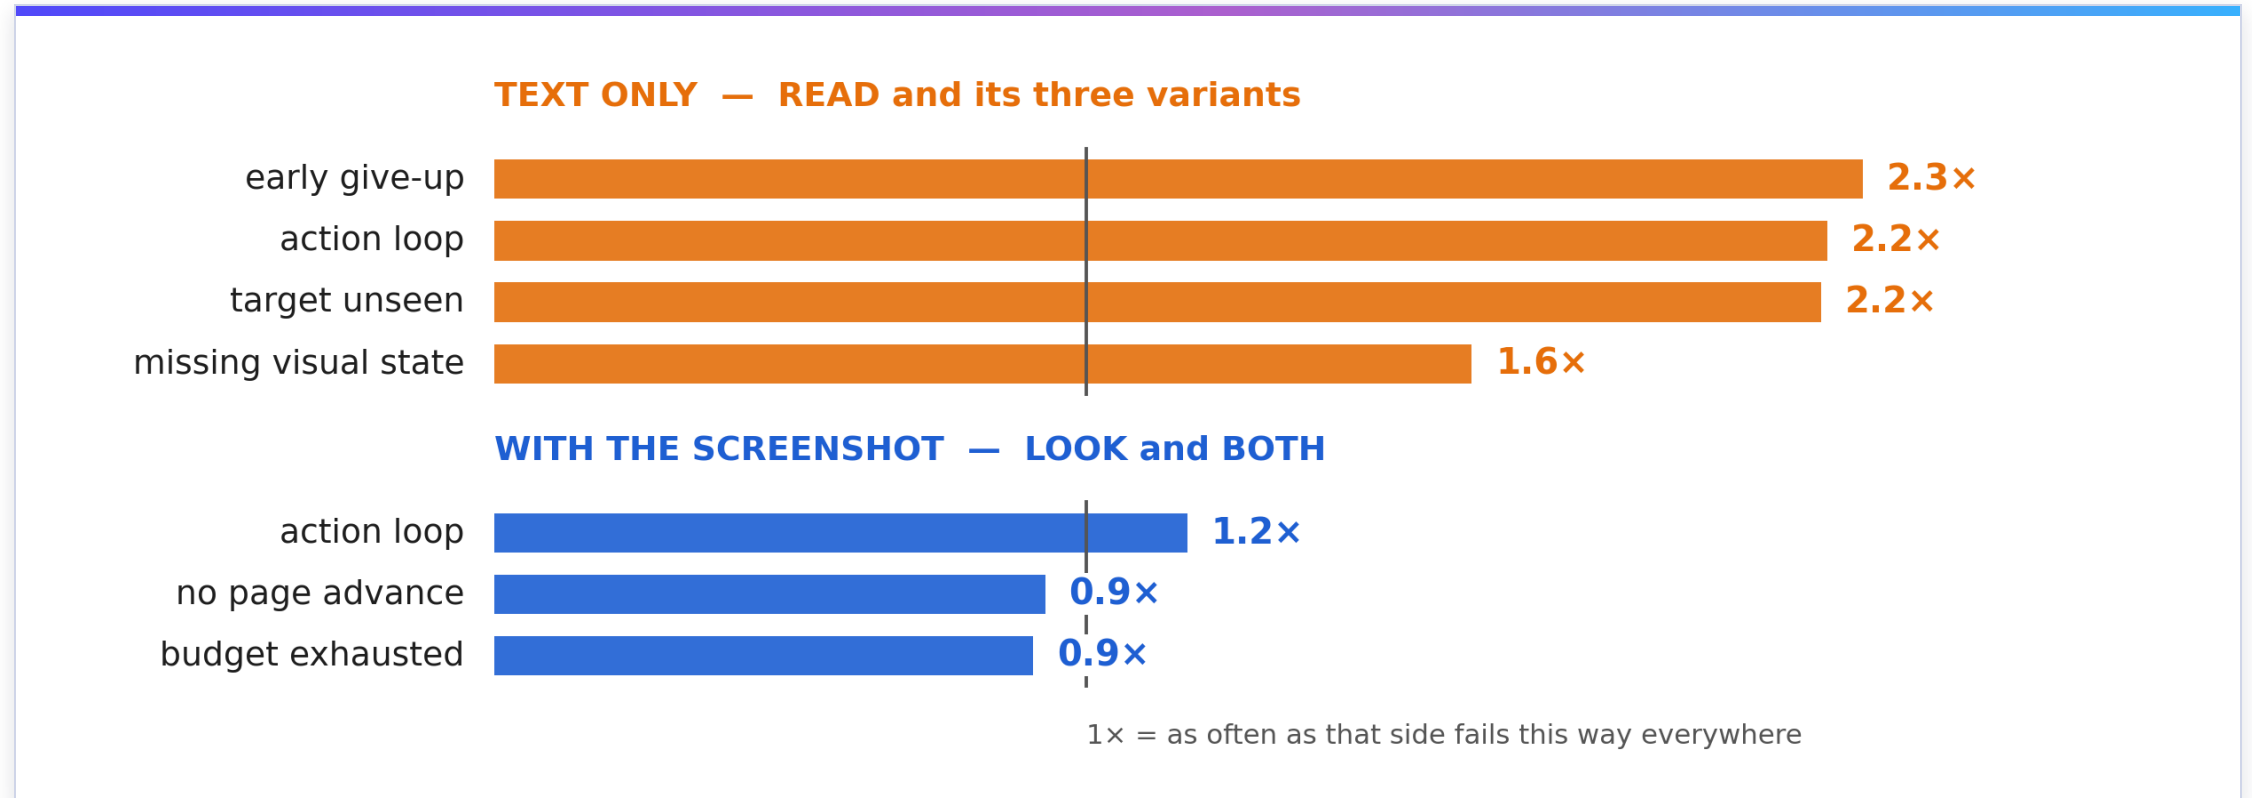
Can it be learned?

■ With only the screenshot, the agent scrolls far more and types far less.



Medians over 8 settings: LOOK scrolls in **about 30%** of its steps, every other view **about 6%**; it types about half as often

- Text-only views fail in ways you can name; views with the screenshot just never get there.



Each side vs its own failures elsewhere · text-only: several over **2x** · LOOK and BOTH: none stands out · six VisualWebArena settings

③ PICKING RIGHT

■ “Claude, find a desk that fits my room.”

If it always picked the right view...




16 more solved. 20% less cost.

Classifieds benchmark · large model · perfect hindsight, not Claude results

④ CAN IT LEARN TO PICK?


■ No learned choice reaches the cheaper-and-no-worse corner.



Learned choices: **0 of 8** settings

Even perfect hindsight: **1 of 8**

■ Routing's scaling law: no wins, no examples.



Fixed-rate illustration, not a fitted law · our estimate where labels are scarce: at least **2–4×** more tasks

NOT YET

■ Improve the agent first, collect reliable examples, then learn when to look.

WHAT IT TAKES

- a stronger agent
- at least **2–4×** more tasks
- reruns that agree

THREE ASKS

- ✓ type a task at the board
- ✓ which view do you use?
- ✓ lend us a stronger agent

Look, read, or both?
It depends on the task.
Nothing we trained knows it yet.



Come to the board — try your own task, live in all three views.

Jiaming Wei · jiaming.wei.25@ucl.ac.uk

REFERENCE – TASK 130, RECORDED STEPPING, 39 S