


A web task, one agent, one browser loop — only the page representation changes



What the thesis finds

1 The headroom is real

Oracle mode choice beats the best fixed mode by **+3.45–16.07pp** (6 VWA cells).
But one rerun of the SAME mode already buys 2.0–7.6pp — so the gap is real, not all of it is representation.

2 The choice is not predictable

With nested CV and two controls, **0/8** cells learn a triage rule that Pareto-beats 'always use the cheapest mode' — a policy that costs nothing to implement.

3 And the reason is structural

A which-mode label only exists once a task is solved, and base SR is **2–27%**.
4 of 6 cells never reach a trainable label supply — resplitting cannot manufacture events.

Scope: VisualWebArena (classifieds, reddit) × three backbones, plus WebArena-reddit × two. Oracle bounds are retrospective, not deployable.