

No observation mode wins everywhere — which is the only reason selection could be worth anything

Task success rate (%), printed absolutely and shaded relative to the best mode within each cell (absolute success spans a factor of ~80 across cells, so an unnormalised map would show only which backbone is stronger). Orange outline = best mode in that cell: 3 different modes win across 8 cells. The rightmost two columns are the external-validation benchmark.

DOM	17.41	14.29	6.25	5.91	1.34	3.94	26.92	16.35
P-prompt	19.64	12.32	6.70	5.42	1.79	0.00	25.96	16.35
P-text	15.62	13.30	7.59	5.91	0.45	1.97	35.58	16.35
P-SoM	15.62	10.84	6.70	5.91	0.89	0.49	25.00	11.54
SoM	27.23	14.78	14.29	7.39	2.23	0.99	22.12	13.46
Vision	25.00	7.39	12.50	2.46	2.23	1.97	19.23	9.62
	cls·B0 n=224	red·B0 n=203	cls·B1 n=224	red·B1 n=203	cls·B2 n=224	red·B2 n=203	wa_red·B0 n=104	wa_red·B1 n=104