



Six corpora, and the places where they are not interchangeable

WA-reddit, WA-shop, WA-shop-admin ship no reference image and no difficulty label. That is a difference in task SPECIFICATION — it does not by itself mean those tasks need less visual grounding.


A Tasks in the corpus, and tasks actually scored




B Annotations the corpus ships with (a deployment has neither)



C What counts as success



D Instruction length (median, with 10-90th percentile)



Source: docs/analysis/benchmark_eda/corpus_eda.json. Panel A's two numbers are why both run-set and scored-set counts appear in this thesis; only the scored set is ever used for a rate. Corpus defects found while building this figure: VWA-clis task 161 visual_difficulty="mediun"; VWA-shop task 179 reasoning_difficulty="hrad" — typos in the benchmark's own annotations, left uncorrected and excluded from any difficulty tally.